Archived example: evaluation task 03 — flipt-io

REAL ARCHIVED ACTIONS / MODEL-CODED LABELS | Intermediate final-test outcomes: UNTESTED

gemini-3.7-flash
54 components | 53 calls | 51,678 run output tokens

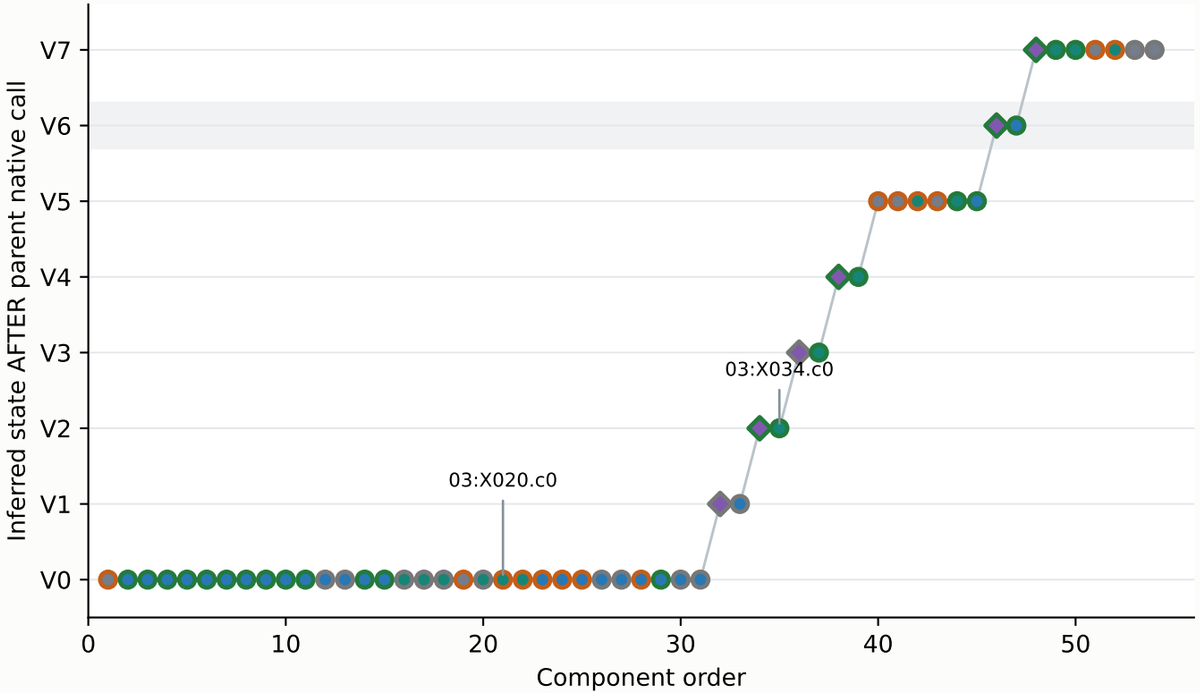

gpt-5.6-luna
30 components | 10 calls | 2,069 run output tokens

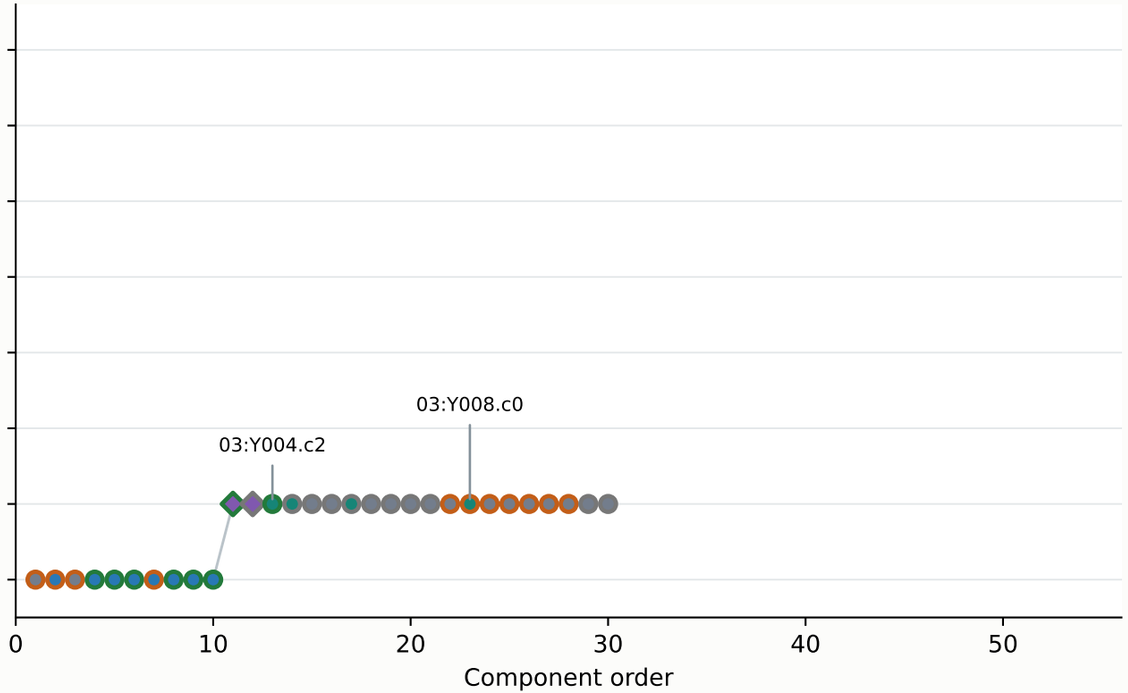

Fill: blue exploration, purple edit, teal verification, gray administration. Diamond = edit.

Border: green direct, orange non-direct, gray uncertain. Labels are archived model judgments; no new annotation.

States follow parent calls; gray bands mark uncertain inferred states. Execution timing inside compound calls is unresolved.

Token amounts above are whole-run totals. Component input/cache/output remain unknown unless actual response usage is recovered.