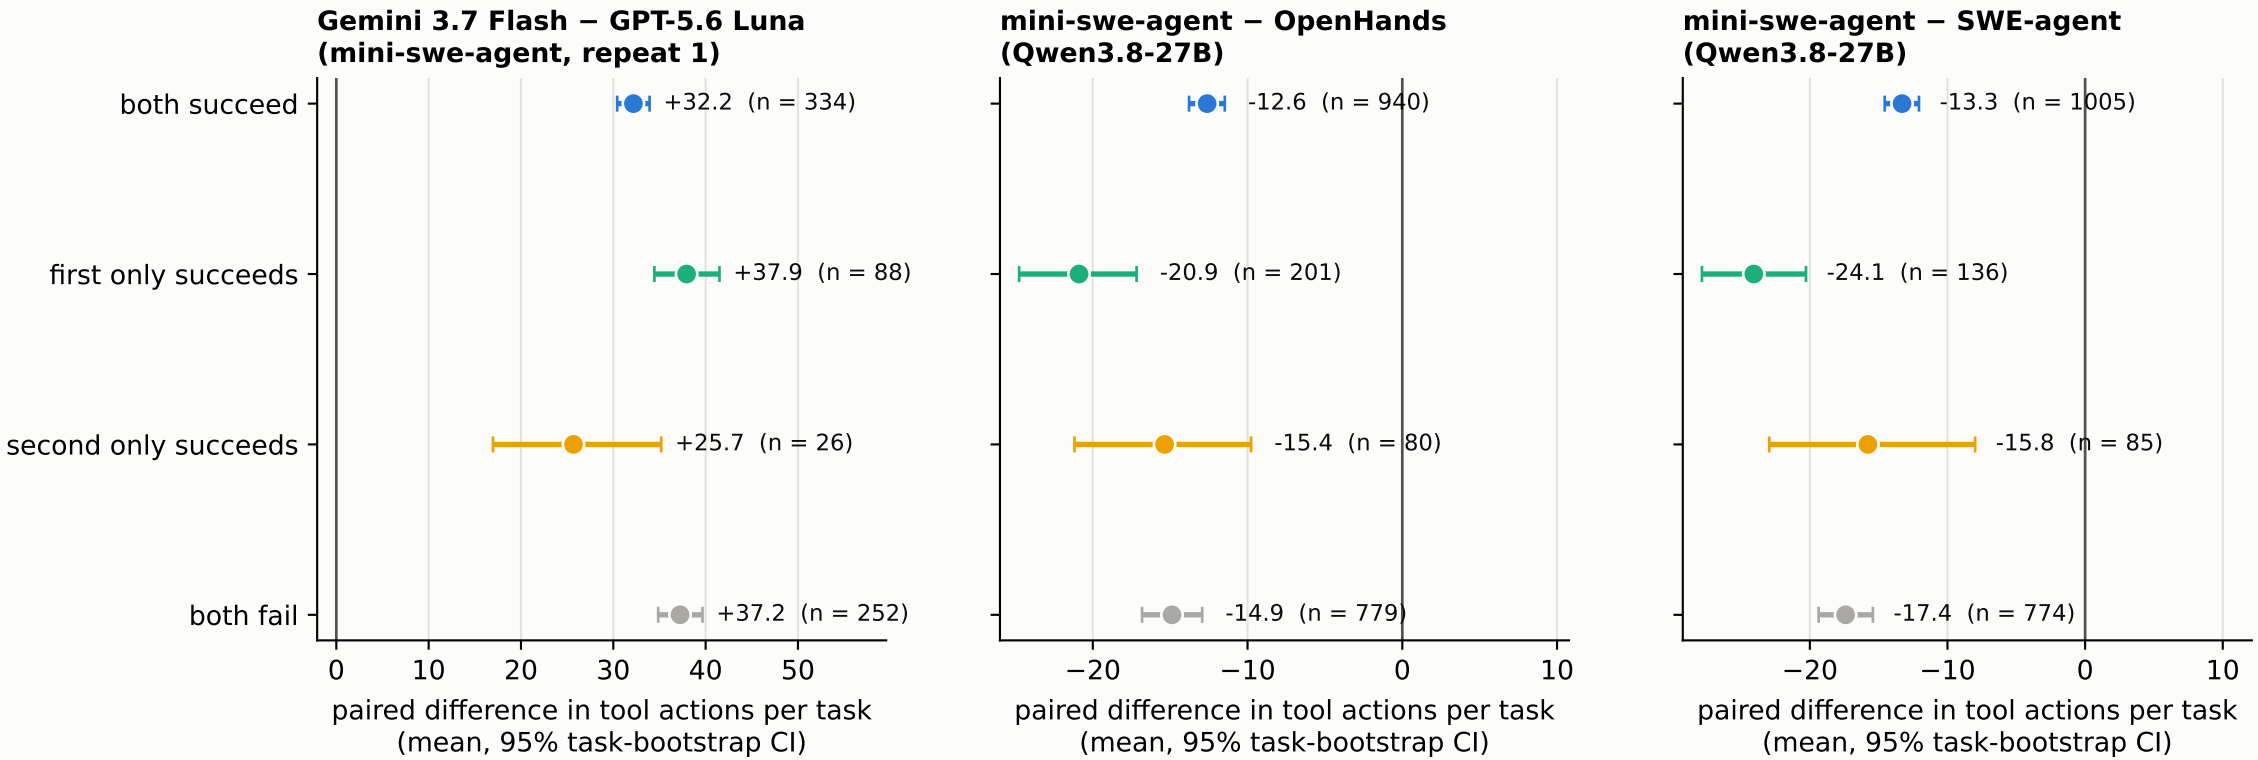
Outcome-matched effort: recorded tool actions per task, paired on task, within each outcome cell
Gemini 3.7 Flash − GPT-5.6 Luna
(mini-swe-agent, repeat 1)
mini-swe-agent − OpenHands
(Qwen3.8-27B)
mini-swe-agent − SWE-agent
(Qwen3.8-27B)
both succeed
first only succeeds
second only succeeds
both fail
+32.2 (n = 334)
+37.9 (n = 88)
+25.7 (n = 26)
+37.2 (n = 252)
-12.6 (n = 940)
-20.9 (n = 201)
-15.4 (n = 80)
-14.9 (n = 779)
-13.3 (n = 1005)
-24.1 (n = 136)
-15.8 (n = 85)
-17.4 (n = 774)
0
10
20
30
40
50
−20
−10
0
10
paired difference in tool actions per task
(mean, 95% task-bootstrap CI)
Saved run tables only; 5,000 task-bootstrap replicates, seed 20261026. Effort is recorded interaction effort conditional on outcome, not redundancy or semantic precision.
Harness call packaging differs (one editor change per call in OpenHands/SWE-agent), so harness differences are not normalized productive work.