Task 03 · Gemini and GPT work progression

Independent horizontal ranges: 54 versus 30 components. Version axis shared; higher state IDs do not mean better code.

Gemini 3.7 Flash | 54 components · 53 calls · 53 responses

Input 3,448,740 (cached 3,200,762)
Output 51,678 | recorded run cost $0.619833

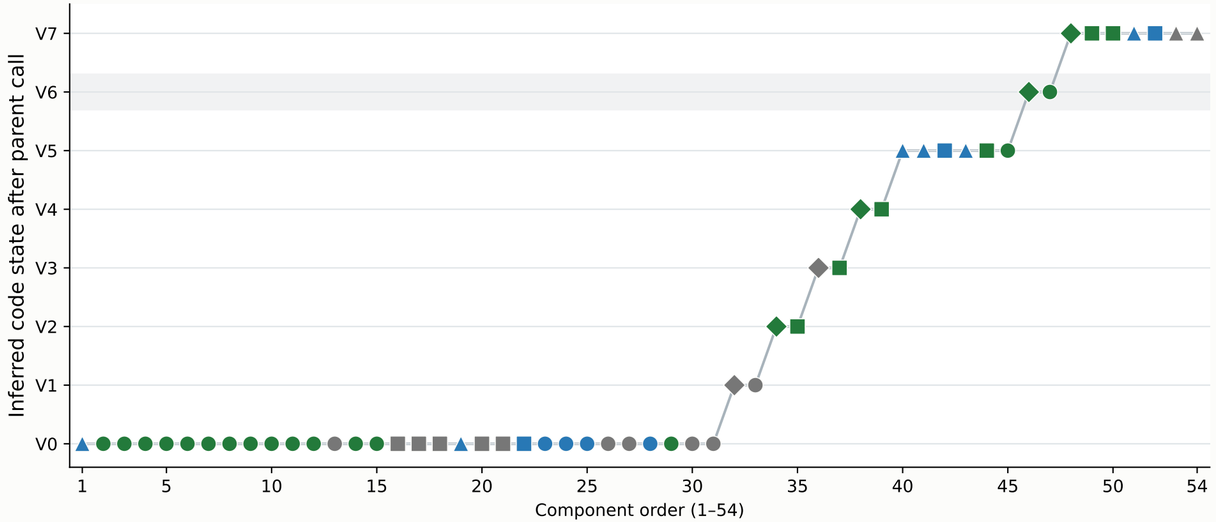

GPT-5.6 Luna | 30 components · 10 calls · 8 responses

Input 100,650 (cached 87,377)
Output 2,069 | recorded run cost $0.006885

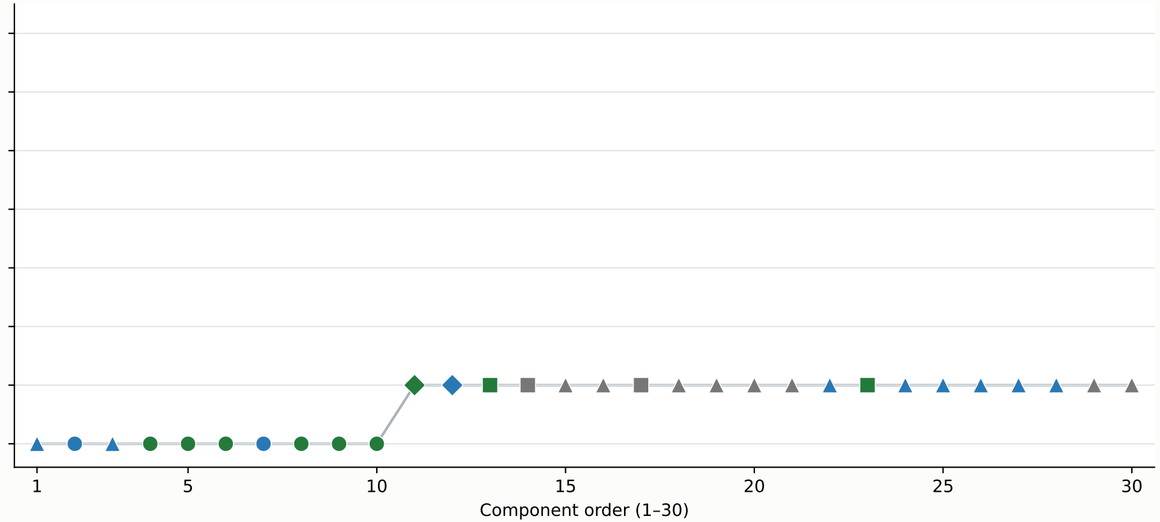

Shape = action type
Exploration · Edit · Verification · Administration

Color = contribution
Direct · Indirect · Uncertain

Indirect is a broad display group combining source indirect, no demonstrated contribution and legacy non-direct. Source annotations and scores are unchanged.

Both panels use the same partial model-review source (5 reviewed cases per run). Compound-call intermediate states are unresolved; independent checkpoints are UNTESTED.

Cached tokens are included in input. USD provenance differs: Gemini source price-map estimate; GPT harness-recorded. Shared-component costs remain allocated estimates.